

Ram Potham

AI control & monitoring researcher. Second author of LinuxArena (Best Paper Award, ICML 2026 AIWILD), used for risk evaluations at Anthropic and the EU AI Office.

ram.potham@gmail.com
github.com/rapturt9
linkedin.com/in/rampotham
[Google Scholar](#)

Research Experience

Redwood Research

Fellow, January 2026 – Current

- Designed and built an automated control-research pipeline: red-team and blue-team agents iterate each protocol against a worst-case attacker.
- Built an automated red-teaming pipeline that turns raw transcripts into sabotage trajectories: the attacks it generates evade elicited frontier monitors **more than 5% of the time**, well above the under-1% baseline for typical model-generated attacks.
- Built LaStraj, a dataset of human-crafted attack trajectories that increases evasion rate **from 1% (model-generated) to 13%** against an Opus 4.6 monitor.
- Second author of LinuxArena (20 environments, 1,671 tasks): ran all evaluations; used for risk evaluations at **Anthropic and the EU AI Office** (see Selected Publications).

Anthropic

Research (Contract), December 2025 – January 2026

- Built AI safety dataset pipelines for Anthropic.

MIT Algorithmic Alignment Lab

AI Safety Researcher, June 2025 – October 2025

- Collaborated with Dylan Hadfield-Menell on character science and model training, using distributed training.

Carnegie Mellon Chimps Lab

AI Researcher, May 2022 – April 2023

- Built WeAudit, a framework for crowd-sourcing failure modes in AI systems.

Selected Publications

LinuxArena: A Control Setting for AI Agents in Live Production Software Environments. 2nd author, Best Paper Award, ICML 2026 Workshop on Agents in the Wild: Safety, Security, and Beyond (Oral Presentation).

Attack Selection in Agentic AI Control Evaluations Meaningfully Decreases Safety. ICML 2026 FAGEN Workshop.

Fair Human-Empowerment-Preserving AI: Desiderata, Design, Behavioral Consequence. 2nd author, 1st IJCAI Workshop on Safe Physical AI, IJCAI-ECAI 2026.

MAEBE: Multi-Agent Emergent Behavior Framework. Equal-contribution author, HICSS Trustworthy AI Track.

Evaluating LLM Agent Adherence to Hierarchical Safety Principles. Sole author, ICML Technical AI Governance Workshop (Oral Presentation).

Industry Experience

Sitewiz (acquired by GAIN)

Founder / CEO, November 2023 – May 2025

- Built and operated autonomous AI agents delivering production client work end-to-end (**10M+ events/day**, team of 3), gaining firsthand experience with agent failure modes, oversight, and reliability at deployment scale.

Education & Recognition

Carnegie Mellon University, School of Computer Science

- Bachelor of Science in Artificial Intelligence, December 2024
- Winner, AI Safety Action Competition, European Network for AI Safety, 2025
- USAMO qualifier, The United States of America Mathematical Olympiad