

Ram Potham

AI control & monitoring researcher, co-author of LinuxArena, used for risk evaluations at both OpenAI and Anthropic.

ram.potham@gmail.com
github.com/rapturt9 ·
linkedin.com/in/rampotham
[Google Scholar](#)

Research Experience

Redwood Research

Fellow, January 2026 – Current

- Co-authored LinuxArena, an agentic coding-agent control arena used for risk evaluations at OpenAI, Anthropic, and in the EU AI Office (see Selected Publications).
- Built an automated control-research pipeline: fleets of AI agents that automate both red-teaming (generating adversarial attacks) and blue-teaming (improving monitors) against control protocols, in an iterative loop.

Anthropic

Research, December 2025 – January 2026

- Built research pipelines for Anthropic, praised as "very valuable" by the lead.

MIT Algorithmic Alignment Lab

AI Safety Researcher, June 2025 – October 2025

- Collaborated with Dylan Hadfield-Menell on character science.

Selected Publications

LinuxArena: A Control Setting for AI Agents in Live Production Software Environments. ICML 2026 Workshop on Agents in the Wild: Safety, Security, and Beyond, Best Paper Award (Oral Presentation).

Attack Selection in Agentic AI Control Evaluations Meaningfully Decreases Safety. ICML 2026 FAGEN Workshop.

Evaluating LLM Agent Adherence to Hierarchical Safety Principles. ICML Technical AI Governance Workshop (Oral Presentation).

Industry Experience

AgentSteer

Founder, March 2026

- Control and monitoring product for coding agents, with users at Coefficient Giving and Redwood Research.
- Found a prompted ~20B open-source model gave the best cost and latency tradeoff as a safeguard: caught prompt injections well, weaker against more sophisticated attacks. Full write-up.

Watertight AI

Technical Staff, November 2025 – December 2025

- Built AI safety datasets and environments for major labs.

Sitewiz (exited to GAIN)

Founder / CEO, November 2023 – May 2025

- Built autonomous AI agents for web development, analytics, and UI/UX automation for leading CRO agencies.

Technical Skills

AI Safety Research AI control and monitoring, safety evaluations, threat modeling

ML Engineering PyTorch, LLM fine-tuning, TensorFlow, numpy, reinforcement learning, transformer architectures

Production Systems Python, AWS, distributed systems, MLOps, autonomous agent deployment

Education & Recognition

Carnegie Mellon University, School of Computer Science

- Bachelor of Science in Artificial Intelligence
- Winner, AI Safety Action Competition, European Network for AI Safety, 2025
- USAMO qualifier, The United States of America Mathematical Olympiad